


DIY Deep Learning for Vision: a Hands-On Tutorial with Caffe



Maximally accurate	Maximally specific
espresso	2.23192
coffee	2.19914
beverage	1.93214
liquid	1.89367
fluid	1.85519



caffe.berkeleyvision.org




github.com/BVLC/caffe

Evan Shelhamer, Jeff Donahue,
Yangqing Jia, Ross Girshick


Look for further
details in the
outline notes



Why Deep Learning? The Unreasonable Effectiveness of Deep Features




Low-level: Pool₁



High-level: FC₆

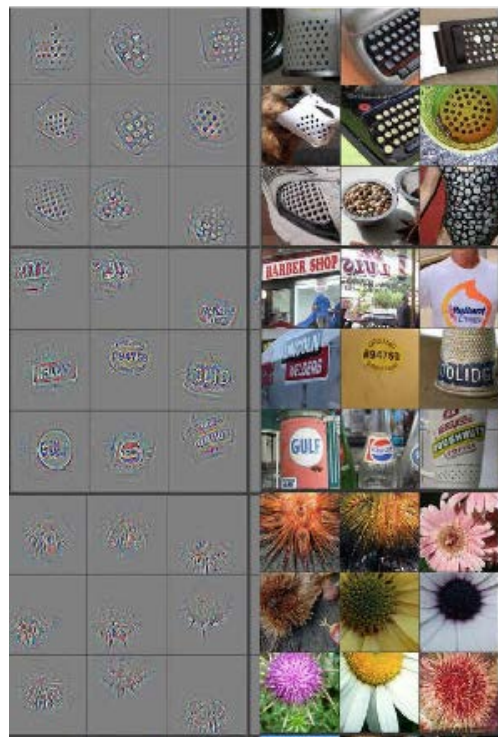
Classes separate in the deep representations and transfer to many tasks.
[DeCAF] [Zeiler-Fergus]

Why Deep Learning? The Unreasonable Effectiveness of Deep Features



Maximal activations of pool₅ units

[R-CNN]



conv₅ DeConv visualization

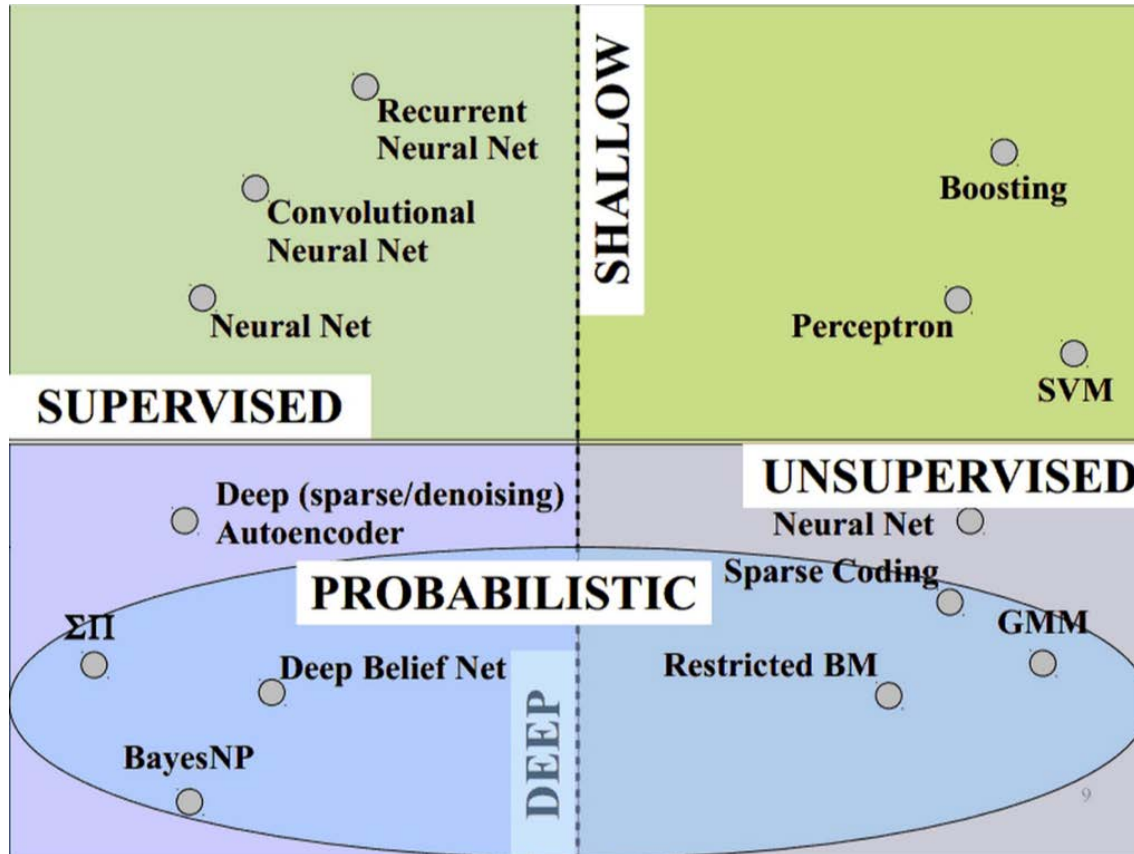
[Zeiler-Fergus]

Rich visual structure of features deep in hierarchy.

What is Deep Learning?

Compositional Models
Learned End-to-End


What is Deep Learning?



Vast space of models!


slide credit Marc'aurelio Ranzato, CVPR '14 tutorial.

Convolutional Neural Nets (CNNs): 1989



LeNet: a layered model composed of convolution and subsampling operations followed by a holistic representation and ultimately a classifier for handwritten digits. [LeNet]

Convolutional Neural Nets (CNNs): 2012



AlexNet: a layered model composed of convolution, subsampling, and further operations followed by a holistic representation and all-in-all a landmark classifier on ILSVRC12. [AlexNet]

- + data
- + gpu
- + non-saturating nonlinearity
- + regularization

Frameworks

- [Torch7](#)
 - NYU
 - scientific computing framework in Lua
 - supported by Facebook
- [Theano/Pylearn2](#)
 - U. Montreal
 - scientific computing framework in Python
 - symbolic computation and automatic differentiation
- [Cuda-Convnet2](#)
 - Alex Krizhevsky
 - Very fast on state-of-the-art GPUs with Multi-GPU parallelism
 - C++ / CUDA library

Framework Comparison

- More alike than different
 - All express deep models
 - All are nicely open-source
 - All include scripting for hacking and prototyping
- No strict winners – experiment and choose the framework that best fits your work
- We like to brew our deep networks with **Caffe**

Why Caffe? In one sip...

- **Expression:** models + optimizations are plaintext schemas, not code.
- **Speed:** for state-of-the-art models and massive data.
- **Modularity:** to extend to new tasks and settings.
- **Openness:** common code and reference models for reproducibility.
- **Community:** joint discussion and development through BSD-2 licensing.

So what is Caffe?

- Pure C++ / CUDA architecture for deep learning
 - command line, Python, MATLAB interfaces
- Fast, well-tested code
- Tools, reference models, demos, and recipes
- Seamless switch between CPU and GPU
 - `Caffe::set_mode(Caffe::GPU);`



Prototype



Training



Deployment

All with essentially the same code!

Caffe is a Community

[project pulse](#)

BVLC / **caffe**

Unwatch

188

★ Unstar

900

Fork

505

August 04 2014 - September 04 2014

Period: 1 month

Overview



69 Active Pull Requests



160 Active Issues

56

Merged Pull Requests

13

Proposed Pull Requests

140

Closed Issues

20

New Issues

Excluding merges, **33 authors** have pushed **64 commits** to master and **520 commits** to all branches. On master, **254 files** have changed and there have been **14,466 additions** and **8,552 deletions**.



1 Release published by 1 person


Published

v0.9999 cold-brew 27 days ago


56 Pull requests merged by 18 people

Reference Models

AlexNet: ImageNet Classification



R-CNN: Regions with CNN features




Caffe offers the

- model definitions
- optimization settings
- pre-trained weights

so you can start right away.


Architectures

DAGs
multi-input
multi-task



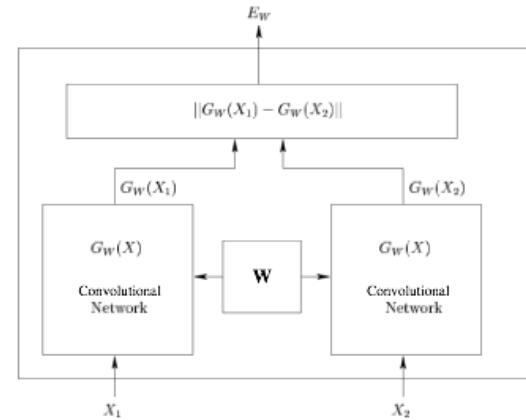
[Karpathy14]

Weight Sharing
Recurrent (RNNs)
Sequences



[Sutskever13]

Siamese Nets
Distances



[Chopra05]

Define your own model from our catalogue
of layers types and start learning.

Brewing by the Numbers...

- Speed with Krizhevsky's 2012 model:
 - K40 / Titan: **2 ms / image**, K20: 2.6ms
 - **40 million images / day**
 - Caffe + cuDNN: **1.17ms / image** on K40
 - 8-core CPU: ~20 ms/image
- **~ 9K** lines of C/C++ code
 - with unit test: ~20k

● C++ 84.2%

● Python 10.5%

● Cuda 3.9%

● Other 1.4%

* Not counting image I/O time. Details at http://caffe.berkeleyvision.org/performance_hardware.html


CAFFE INTRO

Net

- A network is a set of layers connected as a DAG:

```
name: "dummy-net"  
layers { name: "data" ...}  
layers { name: "conv" ...}  
layers { name: "pool" ...}  
... more layers ...  
layers { name: "loss" ...}
```



- Caffe creates and checks the net from the definition.
- Data and derivatives flow through the net as *blobs* – a an array interface



LogReg ↑


LeNet →

ImageNet, Krizhevsky 2012 →



Forward / Backward the essential Net computations

Forward:
inference $f_W(x)$



“espresso”
+ loss

$\nabla f_W(x)$ Backward:
learning


Caffe models are complete machine learning systems for inference and learning. The computation follows from the model definition. Define the model and run.

Layer

```
name: "conv1"
type: CONVOLUTION
bottom: "data"
top: "conv1"
convolution_param {
  num_output: 20
  kernel_size: 5
  stride: 1
  weight_filler {
    type: "xavier"
  }
}
```

name, type, and the
connection structure
(input blobs and
output blobs)

layer-specific
parameters



- Every layer type defines

- **Setup**
- **Forward**
- **Backward**

* Nets + Layers are defined by [protobuf](#) schema

Layer Protocol

Setup: run once for initialization.


Forward: make output given input.

Backward: make gradient of output

- w.r.t. bottom
- w.r.t. parameters (if needed)

Model Composition

The Net forward and backward passes are the composition the layers'.




[Layer Development Checklist](#)

Blob


Blobs are 4-D arrays for storing and communicating information.

- hold data, derivatives, and parameters
- lazily allocate memory
- shuttle between CPU and GPU



Data

Number x K Channel x Height x Width
256 x 3 x 227 x 227 for ImageNet train input



Parameter: Convolution Weight


N Output x K Input x Height x Width
96 x 3 x 11 x 11 for CaffeNet conv1



Parameter: Convolution Bias


96 x 1 x 1 x 1 for CaffeNet conv1

```
name: "conv1"  
type: CONVOLUTION  
bottom: "data"  
top: "conv1"  
... definition ...
```



Blob

Blobs provide a unified memory interface.



Reshape(num, channel, height, width)

- declare dimensions
- make *SyncedMem* -- but only lazily allocate

cpu_data(), mutable_cpu_data()


- host memory for CPU mode

gpu_data(), mutable_gpu_data()

- device memory for GPU mode

{cpu,gpu}_diff(), mutable_{cpu,gpu}_diff()

- derivative counterparts to data methods
- easy access to data + diff in forward / backward



Solving: Training a Net

Optimization like model definition is configuration.

`train_net`: "lenet_train.prototxt"

`base_lr`: 0.01

`momentum`: 0.9

`weight_decay`: 0.0005

`max_iter`: 10000

`snapshot_prefix`: "lenet_snapshot"

`solver_mode`: GPU



All you need to run things on the GPU.

```
> caffe train -solver lenet_solver.prototxt
```

Stochastic Gradient Descent (SGD) + momentum •

Adaptive Gradient (ADAGRAD) • Nesterov's Accelerated Gradient (NAG)

End to End Recipe...

- Convert the data to Caffe-format
 - lmdb, leveldb, hdf5 / .mat, list of images, etc.
- Define the Net
- Configure the Solver
- `caffe train -solver solver.prototxt -gpu 0`
- Examples are your friends
 - `caffe/examples/mnist,cifar10,imagenet`
 - `caffe/build/tools/*`


(Examples)
Logistic Regression
Learn LeNet on MNIST

FINE-TUNING

Fine-tuning


Transferring learned weights to kick-start models

- Take a pre-trained model and fine-tune to new tasks
[DeCAF] [Zeiler-Fergus] [OverFeat]



dog bird invertebrate vehicle good, covering commodity building

Your Task



**Style
Recognition**



**Dogs vs.
Cats**
top 10 in
10 minutes

From ImageNet to Style

- Simply change a few lines in the layer definition

```
layers {  
  name: "data"  
  type: DATA  
  data_param {  
    source: "ilsvrc12_train_leveldb"  
    mean_file: "../../data/ilsvrc12"  
    ...  
  }  
  ...  
}  
...  
layers {  
  name: "fc8"  
  type: INNER_PRODUCT  
  blobs_lr: 1  
  blobs_lr: 2  
  weight_decay: 1  
  weight_decay: 0  
  inner_product_param {  
    num_output: 1000  
    ...  
  }  
}
```

The diagram illustrates the changes needed to adapt a layer definition for style transfer. It shows two code snippets side-by-side, connected by double-headed arrows indicating the transformation.

Input (Left):

```
layers {  
  name: "data"  
  type: DATA  
  data_param {  
    source: "ilsvrc12_train_leveldb"  
    mean_file: "../../data/ilsvrc12"  
    ...  
  }  
  ...  
}  
...  
layers {  
  name: "fc8"  
  type: INNER_PRODUCT  
  blobs_lr: 1  
  blobs_lr: 2  
  weight_decay: 1  
  weight_decay: 0  
  inner_product_param {  
    num_output: 1000  
    ...  
  }  
}
```

Output (Right):

```
layers {  
  name: "data"  
  type: DATA  
  data_param {  
    source: "style_leveldb"  
    mean_file: "../../data/ilsvrc12"  
    ...  
  }  
  ...  
}  
...  
layers {  
  name: "fc8-style"  
  type: INNER_PRODUCT  
  blobs_lr: 1  
  blobs_lr: 2  
  weight_decay: 1  
  weight_decay: 0  
  inner_product_param {  
    num_output: 20  
    ...  
  }  
}
```

Annotations:

- Input:** A different source (points to the `source` change in the `data_param` block).
- Last Layer:** A different classifier (points to the `name` and `num_output` changes in the `inner_product_param` block).
- new name = new params** (points to the `name` change in the `inner_product_param` block).

From ImageNet to Style

```
> caffe train -solver models/finetune_flickr_style/solver.prototxt  
              -weights bvlc_reference_caffenet.caffemodel
```

Under the hood (loosely speaking):

```
net = new Caffe::Net(  
    "style_solver.prototxt");  
net.CopyTrainedNetFrom(  
    pretrained_model);  
solver.Solve(net);
```

Vintage HDR Melancholy Minimal




(Example)

Fine-tuning from ImageNet to Style

LOSS



Loss

What kind of model is this?



Loss

What kind of model is this?



Classification
SOFTMAX_LOSS
HINGE_LOSS
Linear Regression
EUCLIDEAN_LOSS
Attributes / Multiclassification
STANDARD_CROSS_ENTROPY_LOSS
Others...
New Task
NEW_LOSS

Who knows! Need a **loss function**.

Loss

- **Loss function** determines the learning task.
- Given data D , a Net typically minimizes:

$$L(W) = \frac{1}{|D|} \sum_i^{|D|} f_W(X^{(i)}) + \lambda r(W)$$

Data term: error averaged over
instances

Regularization term: penalize large
weights to improve generalization

Loss

- The data error term $f_W(X^{(i)})$ is computed by `Net::Forward`
- Loss is computed as the output of `Layers`
- Pick the loss to suit the task – many different losses for different needs

Softmax Loss Layer

- Multinomial logistic regression: used for predicting a single class of K mutually exclusive classes

```
layers {  
  name: "loss"  
  type: SOFTMAX_LOSS  
  bottom: "pred"  
  bottom: "label"  
  top: "loss"  
}
```

$$\hat{p}_{nk} = \exp(x_{nk}) / [\sum_{k'} \exp(x_{nk'})]$$

$$E = -\frac{1}{N} \sum_{n=1}^N \log(\hat{p}_{n,l_n})$$

Sigmoid Cross-Entropy Loss

- Binary logistic regression: used for predicting K independent probability values in [0, 1]

```
layers {  
  name: "loss"  
  type: SIGMOID_CROSS_ENTROPY_LOSS  
  bottom: "pred"  
  bottom: "label"  
  top: "loss"  
}
```

$$y = (1 + \exp(-x))^{-1}$$

$$E = \frac{1}{N} \sum_{n=1}^N [p_n \log \hat{p}_n + (1 - p_n) \log(1 - \hat{p}_n)]$$

Euclidean Loss

- A loss for regressing to real-valued labels $[-\infty, \infty]$

```
layers {  
  name: "loss"  
  type: EUCLIDEAN_LOSS  
  bottom: "pred"  
  bottom: "label"  
  top: "loss"  
}
```

$$E = \frac{1}{2N} \sum_{n=1}^N \|\hat{y}_n - y_n\|_2^2$$

Multiple loss layers

- Your network can contain as many loss functions as you want
- Reconstruction and Classification:

```
layers {  
  name: "recon-loss"  
  type: EUCLIDEAN_LOSS  
  bottom: "reconstructions"  
  bottom: "data"  
  top: "recon-loss"  
}
```

```
layers {  
  name: "class-loss"  
  type: SOFTMAX_LOSS  
  bottom: "class-preds"  
  bottom: "class-labels"  
  top: "class-loss"  
}
```

$$E = \frac{1}{2N} \sum_{n=1}^N \|\hat{y}_n - y_n\|_2^2 + \frac{1}{N} \sum_{n=1}^N \log(\hat{p}_{n,l_n}),$$

Multiple loss layers

“*_LOSS” layers have a default loss weight of 1

```
layers {  
  name: "loss"  
  type: SOFTMAX_LOSS  
  bottom: "pred"  
  bottom: "label"  
  top: "loss"  
}
```

==

```
layers {  
  name: "loss"  
  type: SOFTMAX_LOSS  
  bottom: "pred"  
  bottom: "label"  
  top: "loss"  
  loss_weight: 1.0  
}
```


Multiple loss layers

- Give each loss its own weight
- E.g. give higher priority to classification error

$$E = \frac{1}{2N} \sum_{n=1}^N \|\hat{y}_n - y_n\|_2^2 + 100^* \frac{1}{N} \sum_{n=1}^N \log(\hat{p}_{n,l_n}),$$

```
layers {  
  name: "recon-loss"  
  type: EUCLIDEAN_LOSS  
  bottom: "reconstructions"  
  bottom: "data"  
  top: "recon-loss"  
}
```

```
layers {  
  name: "class-loss"  
  type: SOFTMAX_LOSS  
  bottom: "class-preds"  
  bottom: "class-labels"  
  top: "class-loss"  
  loss_weight: 100.0  
}
```

Any layer can produce a loss!

- Just add `loss_weight: 1.0` to have a layer's output be incorporated into the loss

$$E = || \text{pred} - \text{label} ||^2 / (2N)$$

```
layers {  
  name: "loss"  
  type: EUCLIDEAN_LOSS  
  bottom: "pred"  
  bottom: "label"  
  top: "euclidean_loss"  
  loss_weight: 1.0  
}
```

==

$$\text{diff} = \text{pred} - \text{label}$$

```
layers {  
  name: "diff"  
  type: ELTWISE  
  bottom: "pred"  
  bottom: "label"  
  top: "diff"  
  eltwise_param {  
    op: SUM  
    coeff: 1  
    coeff: -1  
  }  
}
```

+

$$E = || \text{diff} ||^2 / (2N)$$

```
layers {  
  name: "loss"  
  type: POWER  
  bottom: "diff"  
  top: "euclidean_loss"  
  power_param {  
    power: 2  
  }  
  # = 1/(2N)  
  loss_weight: 0.0078125  
}
```

SOLVER

Solver

- **Solver** optimizes the network weights W to minimize the loss $L(W)$ over the data D

$$L(W) = \frac{1}{|D|} \sum_i^{|D|} f_W \left(X^{(i)} \right) + \lambda r(W)$$

- Coordinates forward / backward, weight updates, and scoring.

Solver

- Computes parameter update ΔW , formed from
 - The stochastic error gradient ∇f_W
 - The regularization gradient $\nabla r(W)$
 - Particulars to each solving method

$$L(W) \approx \frac{1}{N} \sum_i^N f_W (X^{(i)}) + \lambda r(W)$$

SGD Solver

- Stochastic gradient descent, with momentum
- `solver_type`: SGD

$$V_{t+1} = \mu V_t - \alpha \nabla L(W_t)$$

$$W_{t+1} = W_t + V_{t+1}$$

SGD Solver

- “AlexNet” [1] training strategy:
 - Use momentum 0.9
 - Initialize learning rate at 0.01
 - Periodically drop learning rate by a factor of 10
- Just a few lines of Caffe solver specification:

```
base_lr: 0.01 lr_policy:  
"step"  
gamma: 0.1  
stepsize: 100000 max_iter:  
350000  
momentum: 0.9
```

NAG Solver

- Nesterov's accelerated gradient [1]
- `solver_type: NESTEROV`
- Proven to have optimal convergence rate $\mathcal{O}(1/t^2)$ for convex problems

$$V_{t+1} = \mu V_t - \alpha \nabla L(W_t + \mu V_t)$$

$$W_{t+1} = W_t + V_{t+1}$$

AdaGrad Solver

- Adaptive gradient (Duchi et al. [1])
- `solver_type: ADAGRAD`
- Attempts to automatically scale gradients based on historical gradients

$$(W_{t+1})_i = (W_t)_i - \alpha \frac{(\nabla L(W_t))_i}{\sqrt{\sum_{t'=1}^t (\nabla L(W_{t'}))_i^2}}$$

Solver Showdown: MNIST Autoencoder

AdaGrad

SGD

I0901 13:36:30.007884 24952 solver.cpp:232] Iteration 65000, loss = 64.1627

I0901 13:36:30.007922 24952 solver.cpp:251] Iteration 65000, Testing net (#0) # train set

I0901 13:36:33.019305 24952 solver.cpp:289] Test loss: **63.217**

I0901 13:36:33.019356 24952 solver.cpp:302] Test net output #0: cross_entropy_loss = 63.217 (* 1 = 63.217 loss)

I0901 13:36:33.019773 24952 solver.cpp:302] Test net output #1: l2_error = 2.40951

I0901 13:35:20.426187 20072 solver.cpp:232] Iteration 65000, loss = 61.5498

I0901 13:35:20.426218 20072 solver.cpp:251] Iteration 65000, Testing net (#0) # train set

I0901 13:35:22.780092 20072 solver.cpp:289] Test loss: **60.8301**

I0901 13:35:22.780138 20072 solver.cpp:302] Test net output #0: cross_entropy_loss = 60.8301 (* 1 = 60.8301 loss)

I0901 13:35:22.780146 20072 solver.cpp:302] Test net output #1: l2_error = 2.02321

I0901 13:36:52.466069 22488 solver.cpp:232] Iteration 65000, loss = 59.9389

I0901 13:36:52.466099 22488 solver.cpp:251] Iteration 65000, Testing net (#0) # train set

I0901 13:36:55.068370 22488 solver.cpp:289] Test loss: **59.3663**

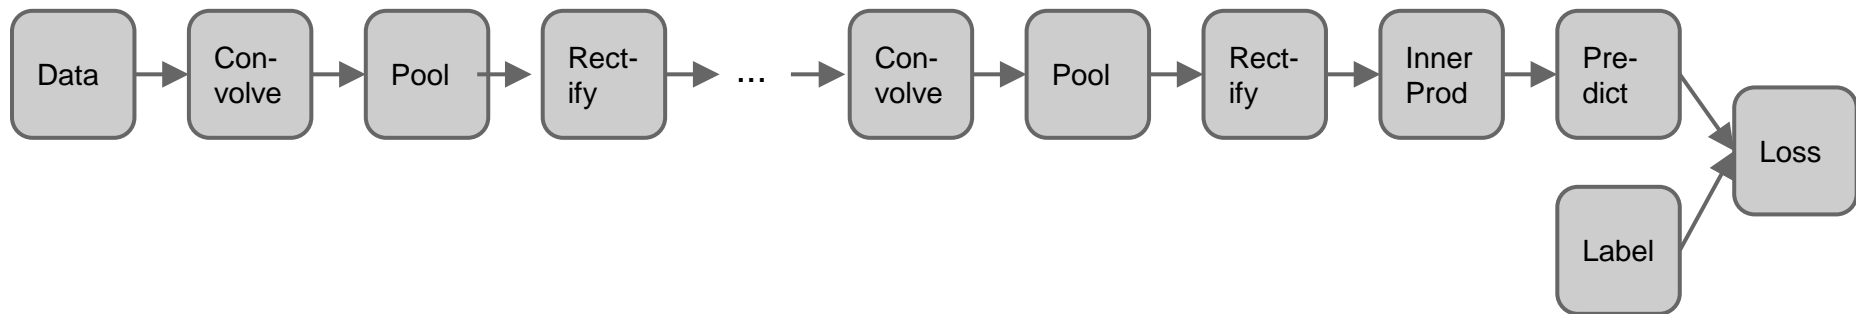
I0901 13:36:55.068410 22488 solver.cpp:302] Test net output #0: cross_entropy_loss = 59.3663 (* 1 = 59.3663 loss)

I0901 13:36:55.068418 22488 solver.cpp:302] Test net output #1: l2_error = 1.79998

DAG

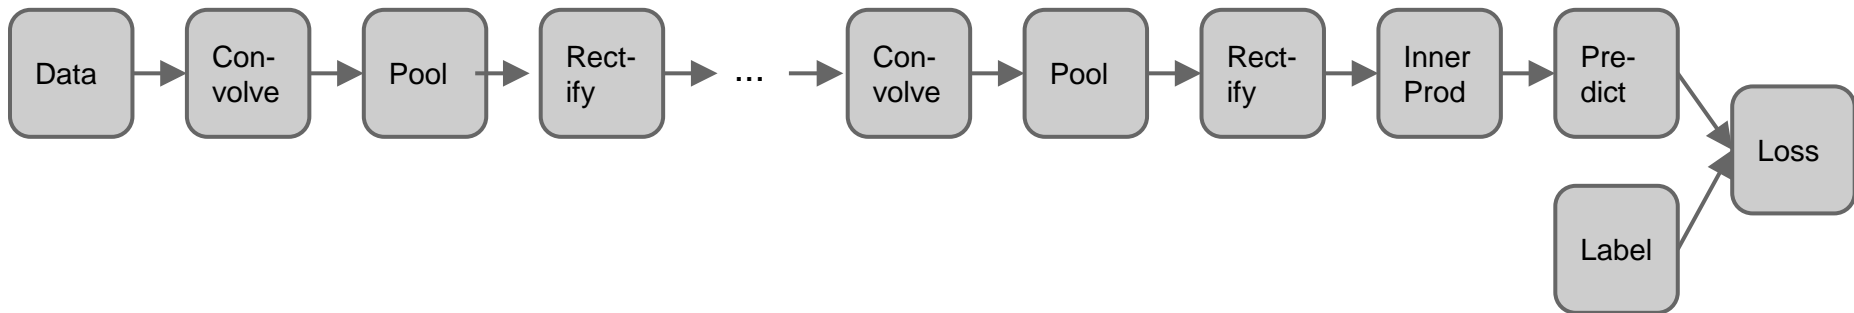
Nets are DAGs

- Modern deep learning approaches to vision have a mostly linear structure

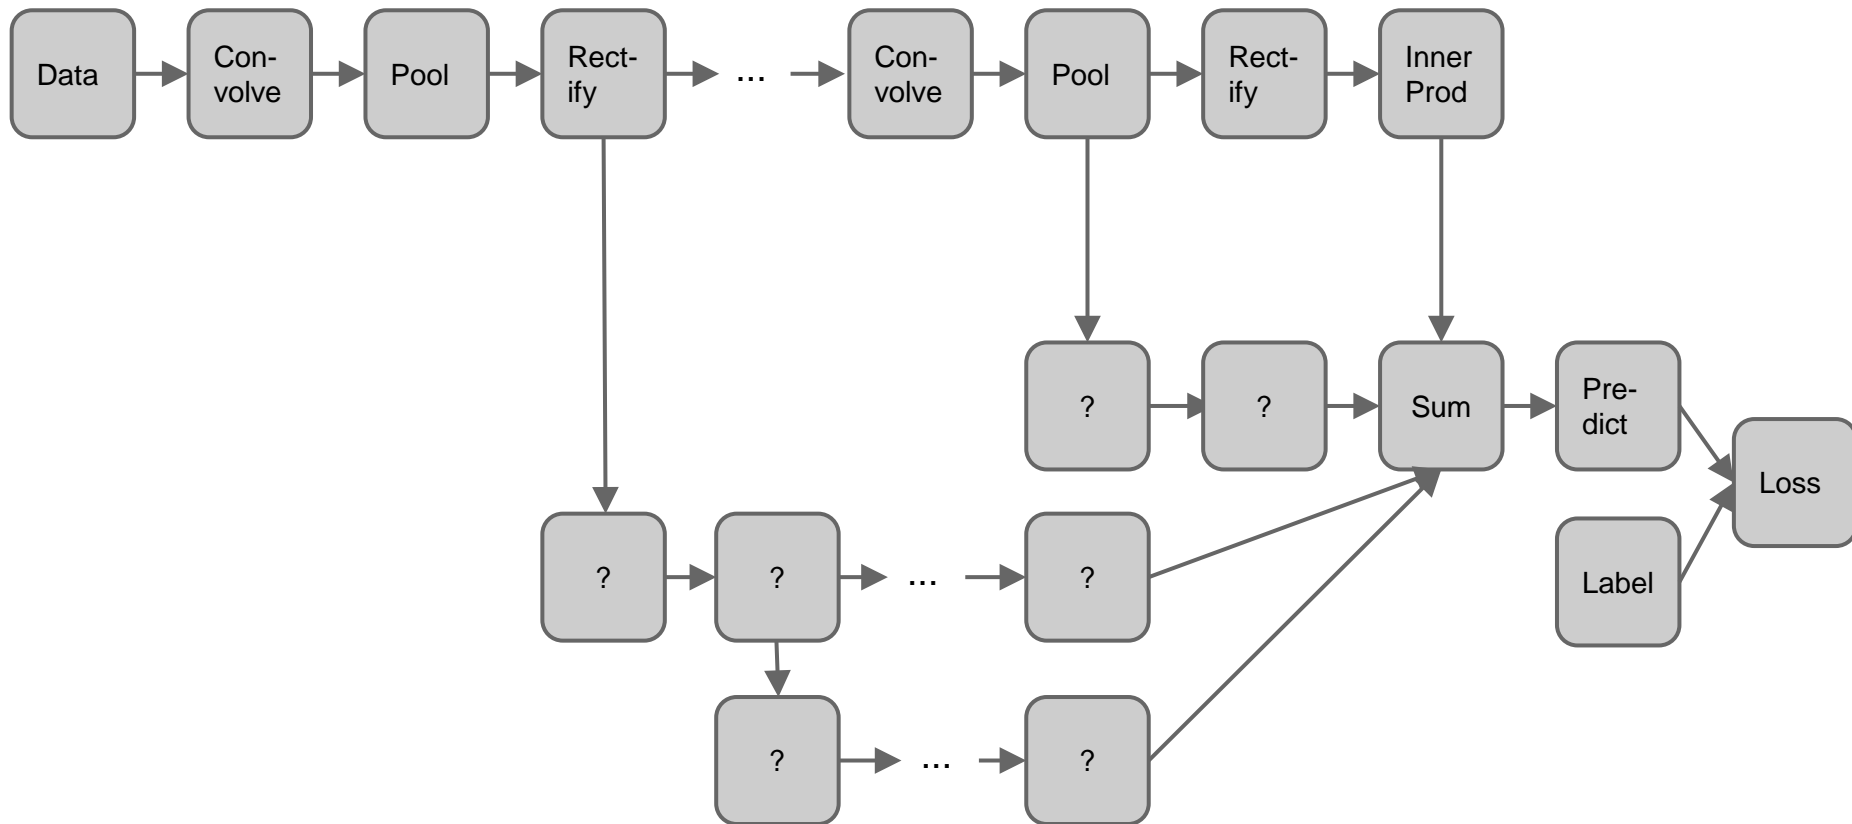


Nets are DAGs

- Modern deep learning approaches to vision have a mostly linear structure
- But Caffe nets can have any arbitrary directed acyclic graph (DAG) structure



Nets are DAGs



WEIGHT SHARING

Weight sharing

- Parameters can be shared and reused across Layers throughout the Net
- Applications:
 - Convolution at multiple scales / pyramids
 - Recurrent Neural Networks (RNNs)
 - Siamese nets for distance learning

Weight sharing

- Just give the parameter blobs explicit names using the `param` field
- Layers specifying the same `param` name will share that parameter, accumulating gradients accordingly

```
layers: {  
  name: 'innerproduct1'  
  type: INNER_PRODUCT  
  inner_product_param {  
    num_output: 10  
    bias_term: false  
    weight_filler {  
      type: 'gaussian'  
      std: 10  
    }  
  }  
  param: 'sharedweights'  
  bottom: 'data'  
  top: 'innerproduct1'  
}  
layers: {  
  name: 'innerproduct2'  
  type: INNER_PRODUCT  
  inner_product_param {  
    num_output: 10  
    bias_term: false  
  }  
  param: 'sharedweights'  
  bottom: 'data'  
  top: 'innerproduct2'  
}
```

EXAMPLES

Share a Sip of Brewed Models

demo.caffe.berkeleyvision.org

demo code open-source and bundled



Maximally accurate

Maximally specific

cat

1.80727

domestic cat

1.74727

feline

1.72787

tabby

0.99133

domestic animal


0.78542

Feature Visualization




The first layer filters, conv1

```
In [8]: # the parameters are a list of [weights, biases]
filters = net.params['conv1'][0].data
vis_square(filters.transpose(0, 2, 3, 1))
```




The first layer output, conv1 (rectified responses of the filters above, first 36 only)

```
In [9]: feat = net.blobs['conv1'].data[4, :36]
vis_square(feat, padval=1)
```




```
In [13]: feat = net.blobs['conv4'].data[4]
vis_square(feat, padval=0.5)
```




The fifth layer output, conv5 (rectified, all 256 channels)

```
In [14]: feat = net.blobs['conv5'].data[4]
vis_square(feat, padval=0.5)
```




The fifth layer after pooling, pool5

```
In [15]: feat = net.blobs['pool5'].data[4]
vis_square(feat, padval=1)
```




```
In [16]: feat = net.blobs['fc6'].data[4]
plt.subplot(2, 1, 1)
plt.plot(feat.flat)
plt.subplot(2, 1, 2)
plt.hist(feat.flat(feat.flat > 0), bins=100)
```




The second fully connected layer, fc7 (rectified)

```
In [17]: feat = net.blobs['fc7'].data[4]
plt.subplot(2, 1, 1)
plt.plot(feat.flat)
plt.subplot(2, 1, 2)
plt.hist(feat.flat(feat.flat > 0), bins=100)
```






```
In [18]: feat = net.blobs['prob'].data[4]
plt.plot(feat.flat)
```

```
Out[18]: <matplotlib.lines.Line2D at 0x12b60710>
```



Fully-convolutional Models



Transform fixed-input models into any-size models by translating inner products to convolutions.

The computation exploits a natural efficiency of convolutional neural network (CNN) structure by dynamic programming in the forward pass from shallow to deep layers and analogously in backward.

[Net surgery in Caffe](#)

how to transform models:

- make fully-convolutional
- set custom weights

Object Detection

R-CNN: Regions with Convolutional Neural Networks


<http://nbviewer.ipython.org/github/BVLC/caffe/blob/master/examples/detection.ipynb>

Full R-CNN scripts available at


<https://github.com/rbgirshick/rcnn>

Ross Girshick et al.

Rich feature hierarchies for accurate object detection and semantic segmentation. CVPR14.



Thanks to the Caffe crew



Yangqing Jia, Evan Shelhamer, Jeff Donahue, Sergey Karayev
Jonathan Long, Ross Girshick, Sergio Guadarrama

and our [open source contributors](#)!



...plus the cold-brew

Acknowledgements



Thank you to the Berkeley Vision and Learning Center Sponsors.



Thank you to NVIDIA
for collaboration on cuDNN
and GPU donation.



Thank you to our 30+
open source contributors
and vibrant community.

References

- [DeCAF] J. Donahue, Y. Jia, O. Vinyals, J. Hoffman, N. Zhang, E. Tzeng, and T. Darrell. Decaf: A deep convolutional activation feature for generic visual recognition. ICML, 2014.
- [R-CNN] R. Girshick, J. Donahue, T. Darrell, and J. Malik. Rich feature hierarchies for accurate object detection and semantic segmentation. CVPR, 2014.
- [Zeiler-Fergus] M. Zeiler and R. Fergus. Visualizing and understanding convolutional networks. ECCV, 2014.
- [LeNet] Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. IEEE, 1998.
- [AlexNet] A. Krizhevsky, I. Sutskever, and G. Hinton. Imagenet classification with deep convolutional neural networks. NIPS, 2012.
- [OverFeat] P. Sermanet, D. Eigen, X. Zhang, M. Mathieu, R. Fergus, and Y. LeCun. Overfeat: Integrated recognition, localization and detection using convolutional networks. ICLR, 2014.
- [Image-Style] [S. Karayev](#), [M. Trentacoste](#), [H. Han](#), [A. Agarwala](#), [T. Darrell](#), [A. Hertzmann](#), [H. Winnemoeller](#). Recognizing Image Style. BMVC, 2014.
- [Karpathy14] A. Karpathy, G. Toderici, S. Shetty, T. Leung, R. Sukthankar, and L. Fei-Fei. Large-scale video classification with convolutional neural networks. CVPR, 2014.
- [Sutskever13] I. Sutskever. Training Recurrent Neural Networks. PhD thesis, University of Toronto, 2013.
- [Chopra05] S. Chopra, R. Hadsell, and Y. LeCun. Learning a similarity metric discriminatively, with application to face verification. CVPR, 2005.